

Introduction to Embedded Systems

Edward A. Lee

UC Berkeley
EECS 149/249A
Fall 2016

© 2008-2016: E. A. Lee, A. L. Sangiovanni-Vincentelli, S. A. Seshia.
All rights reserved.

Chapter 9: Memory Architectures

Role of Memory in Embedded Systems

Traditional roles: Storage and Communication for Programs

Communication with Sensors and Actuators

Often much more constrained than in general-purpose computing

- Size, power, reliability, etc.

Can be important for programmers to understand these constraints

Memory Architecture: Issues

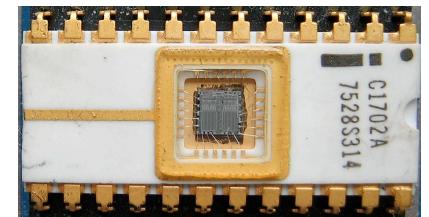
- Types of memory
 - volatile vs. non-volatile, SRAM vs. DRAM
- Memory maps
 - Harvard architecture
 - Memory-mapped I/O
- Memory organization
 - statically allocated
 - stacks
 - heaps (allocation, fragmentation, garbage collection)
- The memory model of C
- Memory hierarchies
 - scratchpads, caches, virtual memory)
- Memory protection
 - segmented spaces

These issues loom larger in embedded systems than in general-purpose computing.

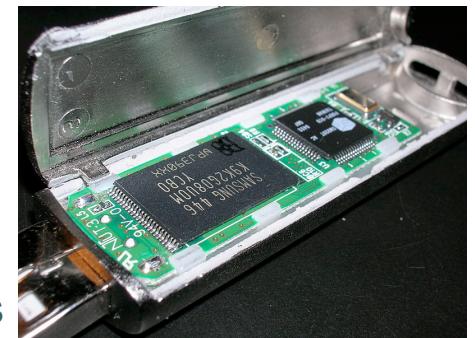
Non-Volatile Memory

Preserves contents when power is off

- **EPROM:** erasable programmable read only memory
 - Invented by Dov Frohman of Intel in 1971
 - Erase by exposing the chip to strong UV light
- **EEPROM:** electrically erasable programmable read-only memory
 - Invented by George Perlegos at Intel in 1978
- **Flash memory**
 - Invented by Dr. Fujio Masuoka at Toshiba around 1980
 - Erased a “block” at a time
 - Limited number of program/erase cycles (~ 100,000)
 - Controllers can get quite complex
- **Disk drives**
 - Not as well suited for embedded systems



USB Drive



Images from the Wikimedia Commons

Volatile Memory

Loses contents when power is off.

- SRAM: static random-access memory
 - Fast, deterministic access time
 - But more power hungry and less dense than DRAM
 - Used for caches, scratchpads, and small embedded memories
- DRAM: dynamic random-access memory
 - Slower than SRAM
 - Access time depends on the sequence of addresses
 - Denser than SRAM (higher capacity)
 - Requires periodic refresh (typically every 64msec)
 - Typically used for main memory
- Boot loader
 - On power up, transfers data from non-volatile to volatile memory.

Example:

Die of a
STM32F103VGT6
ARM Cortex-M3
microcontroller with
1 megabyte flash
memory by
STMicroelectronics.

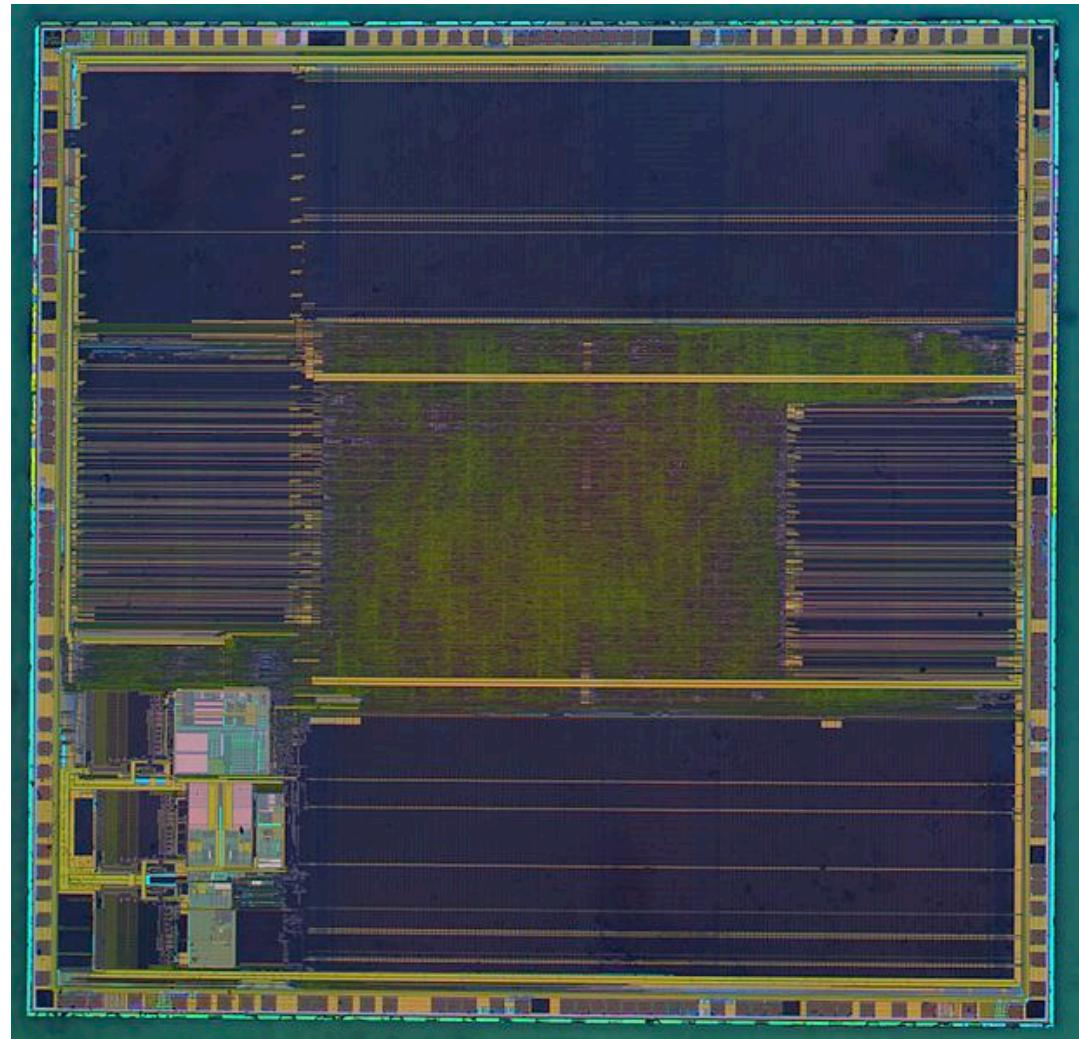



Image from Wikimedia Commons

Memory Map of an ARM Cortex™ - M3 architecture

Defines the
mapping of
addresses to
physical memory.

Note that this does
not define how
much physical
memory there is!



Another Example: AVR



The AVR is an 8-bit single chip microcontroller first developed by Atmel in 1997. The AVR was one of the first microcontroller families to use on-chip flash memory for program storage. It has a modified Harvard architecture.¹

AVR was conceived by two students at the Norwegian Institute of Technology (NTH) Alf-Egil Bogen and Vegard Wollan, who approached Atmel in Silicon Valley to produce it.

¹ A Harvard architecture uses separate memory spaces for program and data. It originated with the Harvard Mark I relay-based computer (used during World War II), which stored the program on punched tape (24 bits wide) and the data in electro-mechanical counters.

A Use of AVR: Arduino

Arduino is a family of open-source hardware boards built around either 8-bit AVR processors or 32-bit ARM processors.

Example:
Atmel AVR
Atmega328
28-pin DIP on an
Arduino Duemilanove
board




Image from Wikimedia Commons EECS 149/249A, UC Berkeley: 9

Open-Source Hardware and the maker movement



Massimo Banzi, founder of the Arduino project at Ivrea, Italy, and Limor Fried, owner and founder of Adafruit, showing one of the first board Arduino Uno from the production lines of Adafruit.

[<http://www.open-electronics.org>]

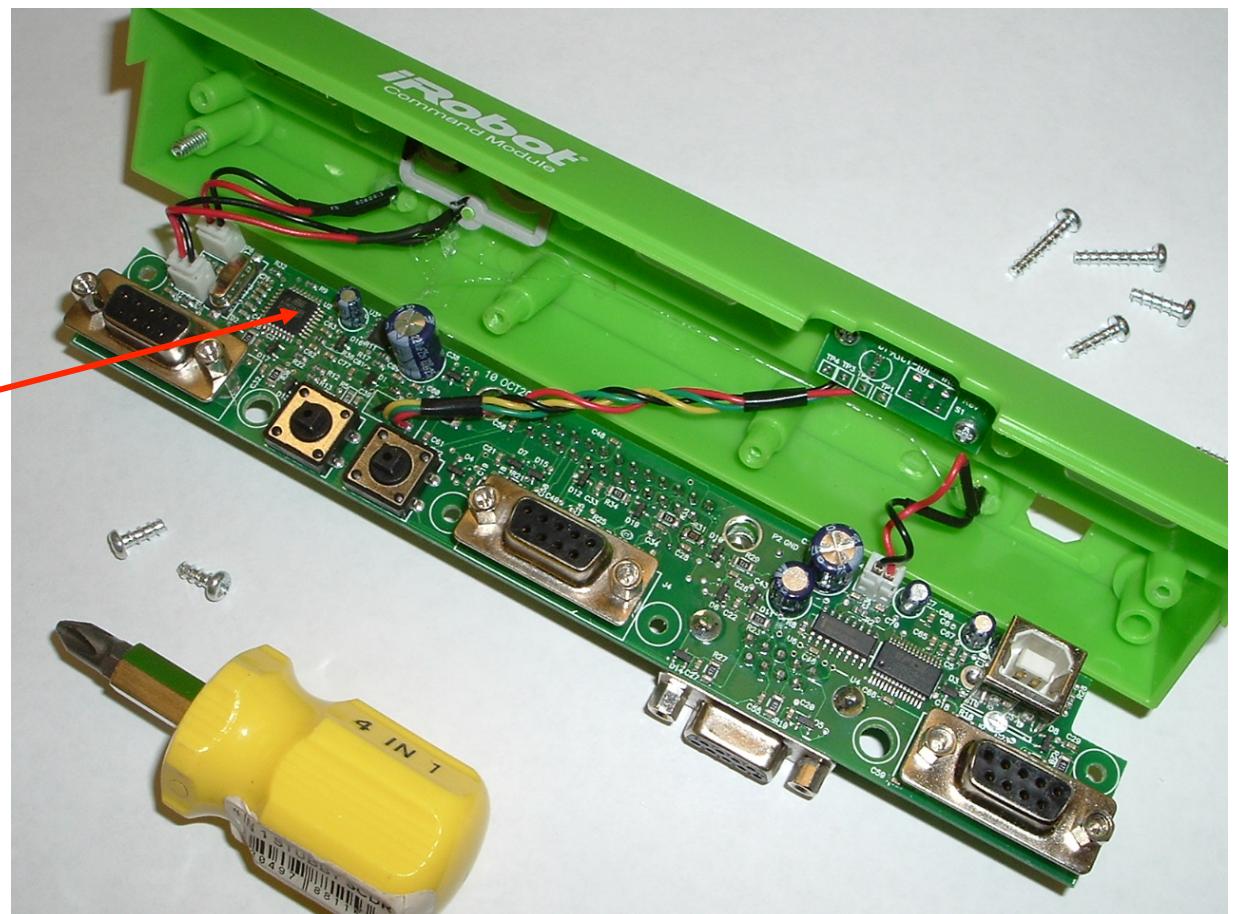

EECS 149/249A, UC Berkeley: 10

Another example use
of an AVR processor

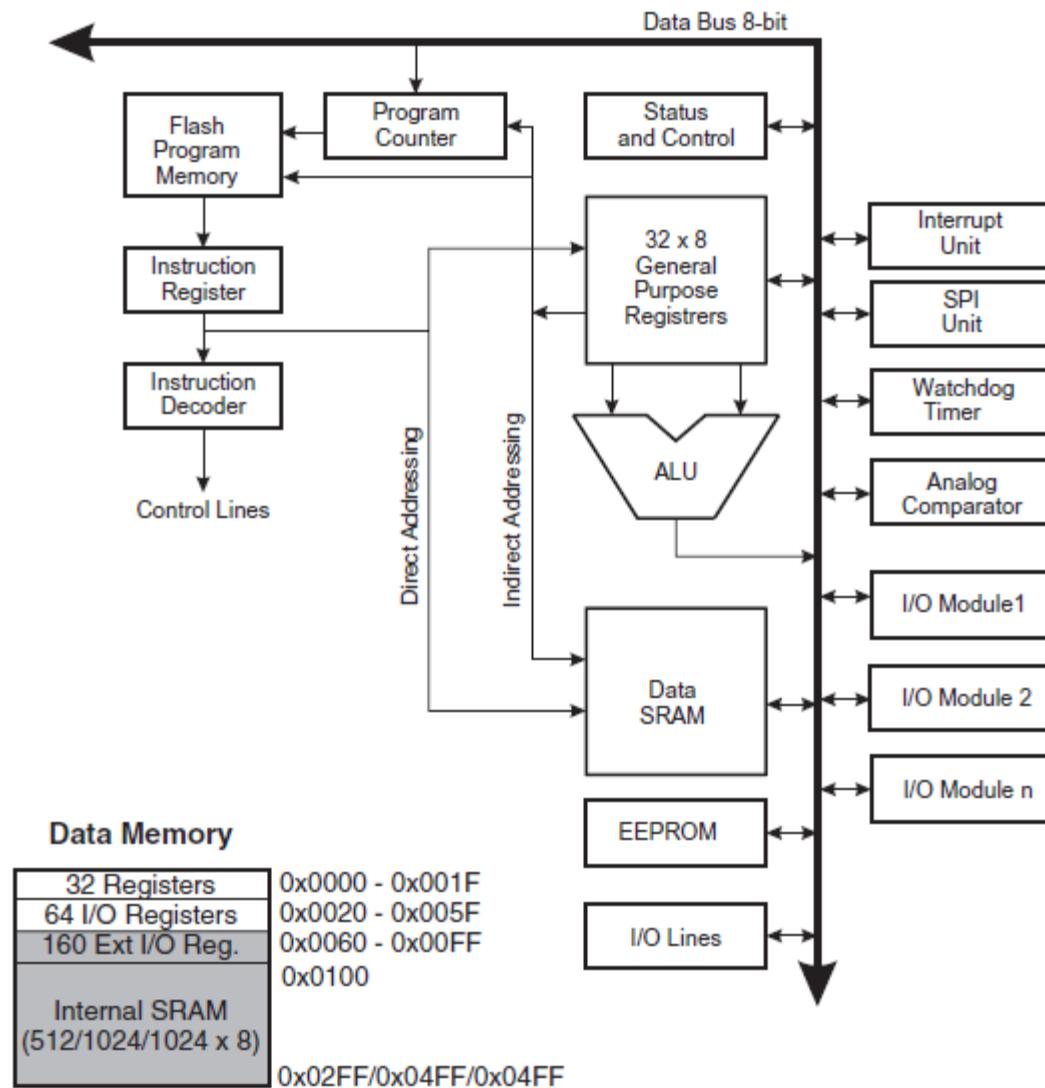
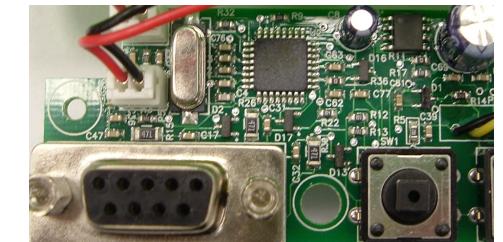


The
iRobot Create
Command Module

Atmel ATmega 168
Microcontroller



ATMega 168: An 8-bit microcontroller with 16-bit addresses




AVR microcontroller architecture used in iRobot command module.

Why is it called an 8-bit microcontroller?

ATMega168 Memory Architecture

An 8-bit microcontroller with 16-bit addresses




iRobot command module has 16K bytes flash memory (14,336 available for the user program. Includes interrupt vectors and boot loader.)

Data Memory

32 Registers	0x0000 - 0x001F
64 I/O Registers	0x0020 - 0x005F
160 Ext I/O Reg.	0x0060 - 0x00FF
Internal SRAM (512/1024/1024 x 8)	0x0100

0x02FF/0x04FF/0x04FF



The “8-bit data” is why this is called an “8-bit microcontroller.”

Additional I/O on the command module:

- Two 8-bit timer/counters
- One 16-bit timer/counter
- 6 PWM channels
- 8-channel, 10-bit ADC
- One serial UART
- 2-wire serial interface

Source: ATmega168 Reference Manual

EECS 149/249A, UC Berkeley: 13

Questions to test your understanding

1. What is the difference between an 8-bit microcontroller and a 32-bit microcontroller?
2. Why use volatile memory? Why not always use non-volatile memory?

Memory Organization for Programs

- Statically-allocated memory
 - Compiler chooses the address at which to store a variable.
- Stack
 - Dynamically allocated memory with a Last-in, First-out (LIFO) strategy
- Heap
 - Dynamically allocated memory

Statically-Allocated Memory in C

```
char x;  
int main(void) {  
    x = 0x20;  
    ...  
}
```

Compiler chooses what address to use for x, and the variable is accessible across procedures. The variable's lifetime is the total duration of the program execution.

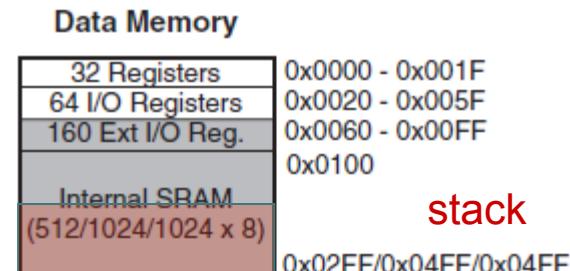
Statically-Allocated Memory with Limited Scope

```
void foo(void)  {
    static char x;
    x = 0x20;
    ...
}
```

Compiler chooses what address to use for x, but the variable is meant to be accessible only in foo(). The variable's lifetime is the total duration of the program execution (values persist across calls to foo()).

Variables on the Stack ("automatic variables")

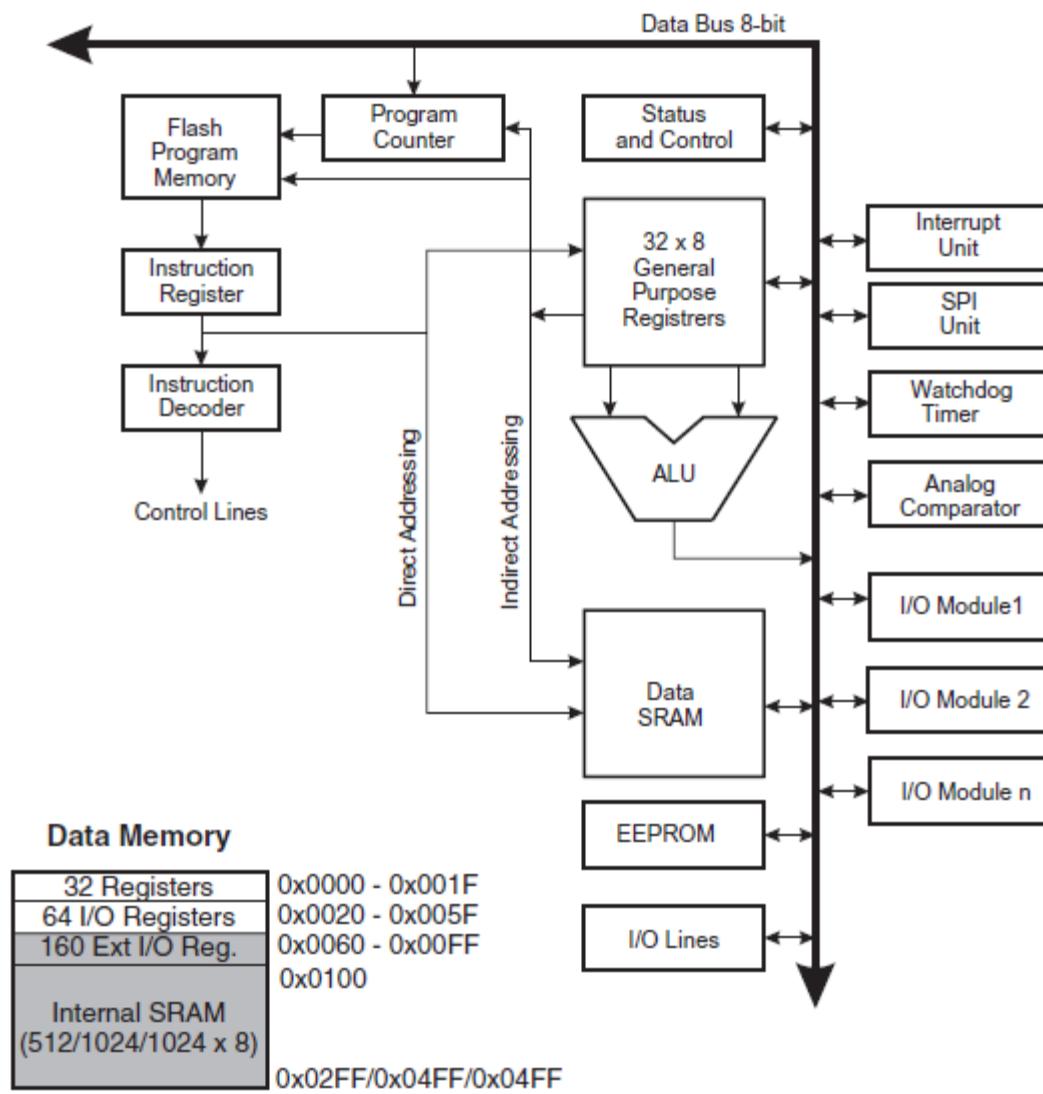
```
void foo(void) {  
    char x;  
    x = 0x20;  
    ...  
}
```



As nested procedures get called, the stack pointer moves to lower memory addresses. When these procedures return, the pointer moves up.

When the procedure is called, x is assigned an address on the stack (by decrementing the stack pointer). When the procedure returns, the memory is freed (by incrementing the stack pointer). The variable persists only for the duration of the call to foo().

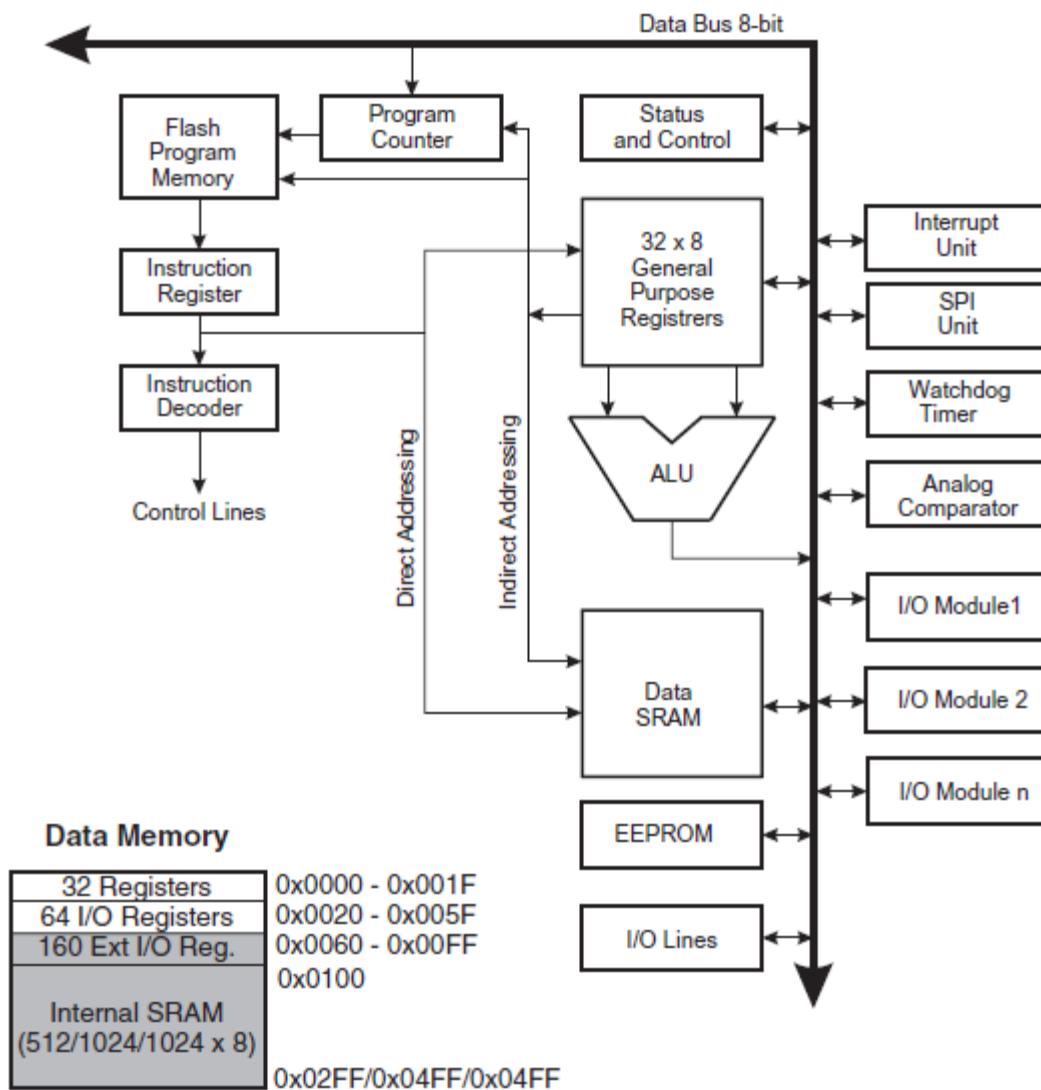
Question 1



What is meant by the following C code:

```
char x;  
void foo(void) {  
    x = 0x20;  
    ...  
}
```

Answer 1




What is meant by the following C code:

```
char x;  
void foo(void) {  
    x = 0x20;  
    ...  
}
```

An 8-bit quantity (hex 0x20) is stored at an address in statically allocated memory in internal RAM determined by the compiler.

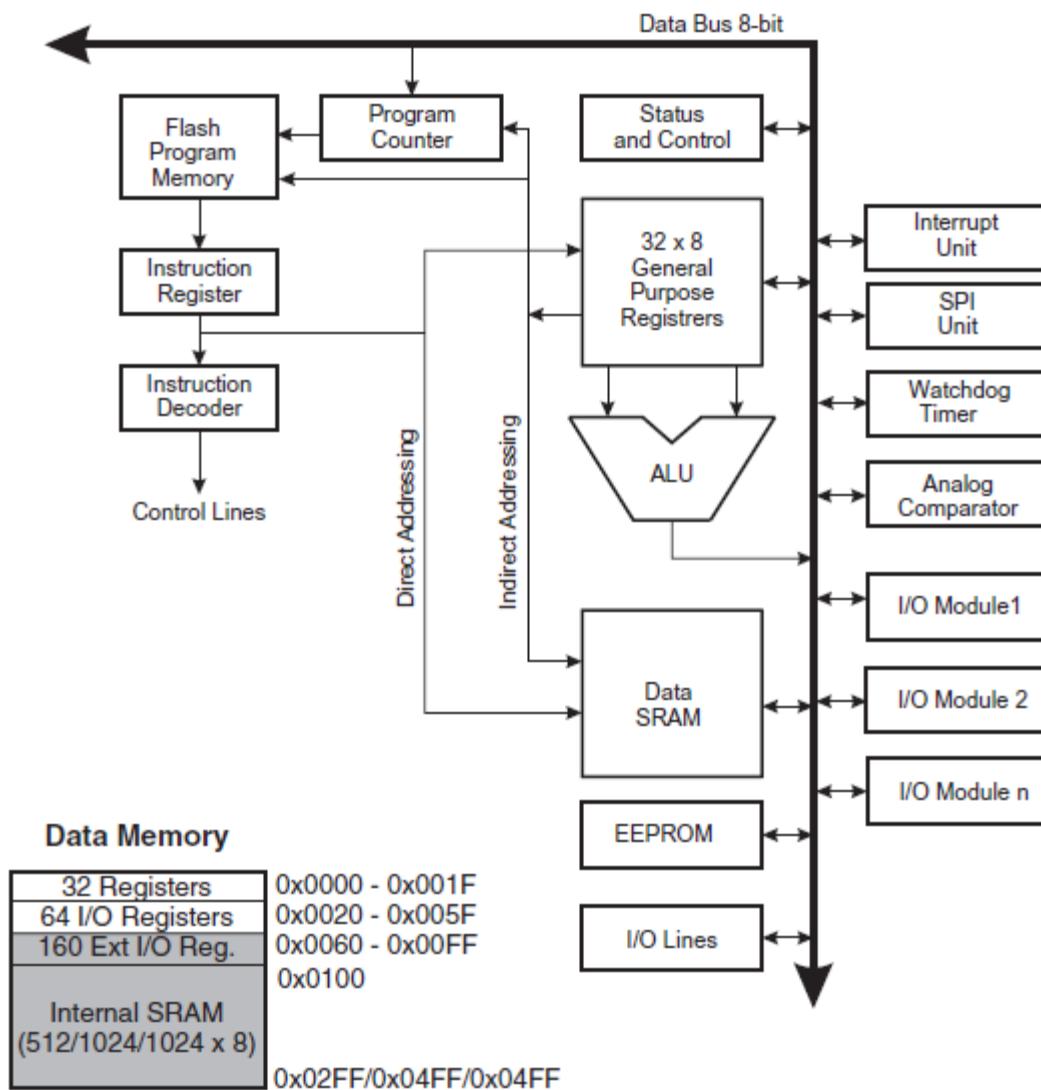
Question 2



What is meant by the following C code:

```
char *x;  
void foo(void) {  
    x = 0x20;  
    ...  
}
```

Answer 2

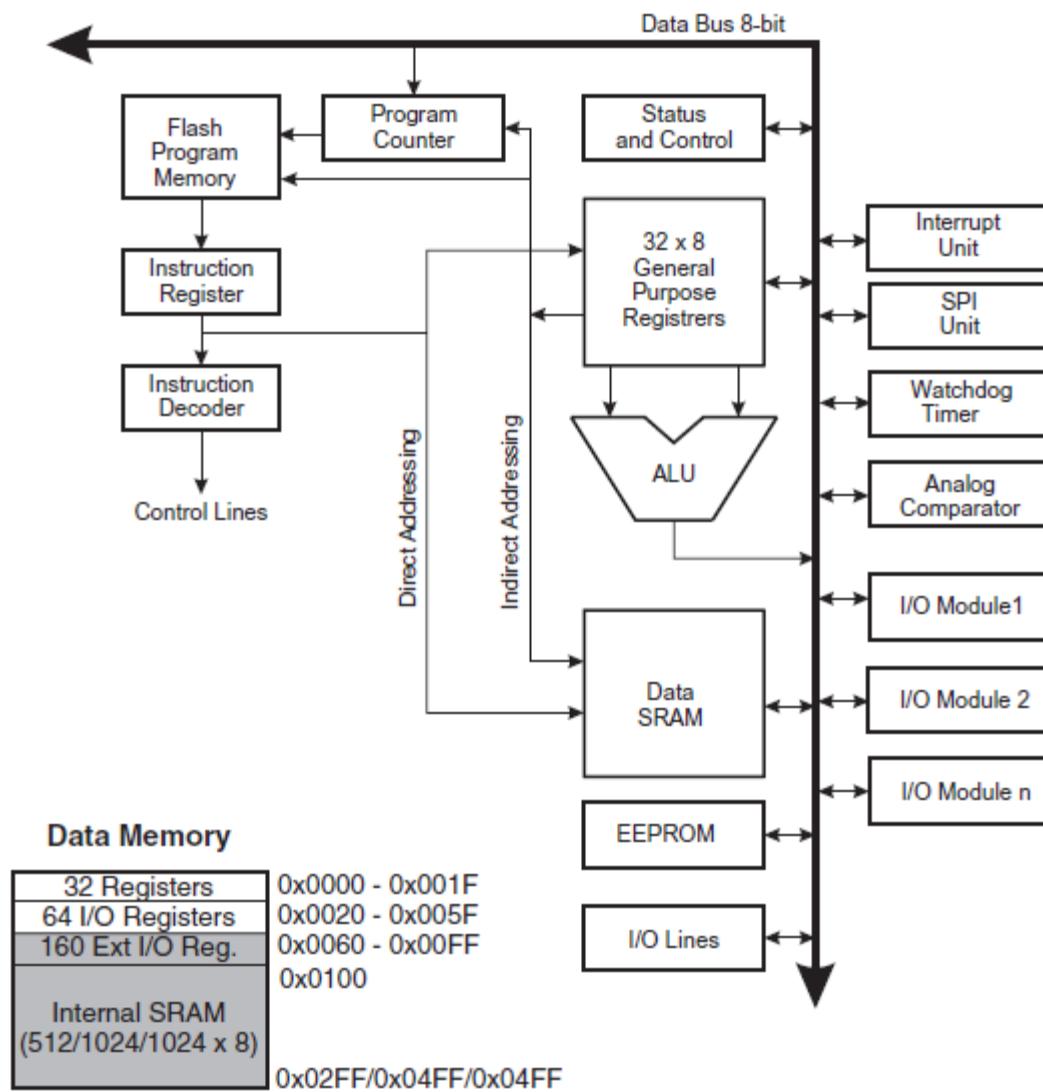


What is meant by the following C code:

```
char *x;  
void foo(void) {  
    x = 0x20;  
    ...  
}
```

An 16-bit quantity (hex 0x0020) is stored at an address in statically allocated memory in internal RAM determined by the compiler.

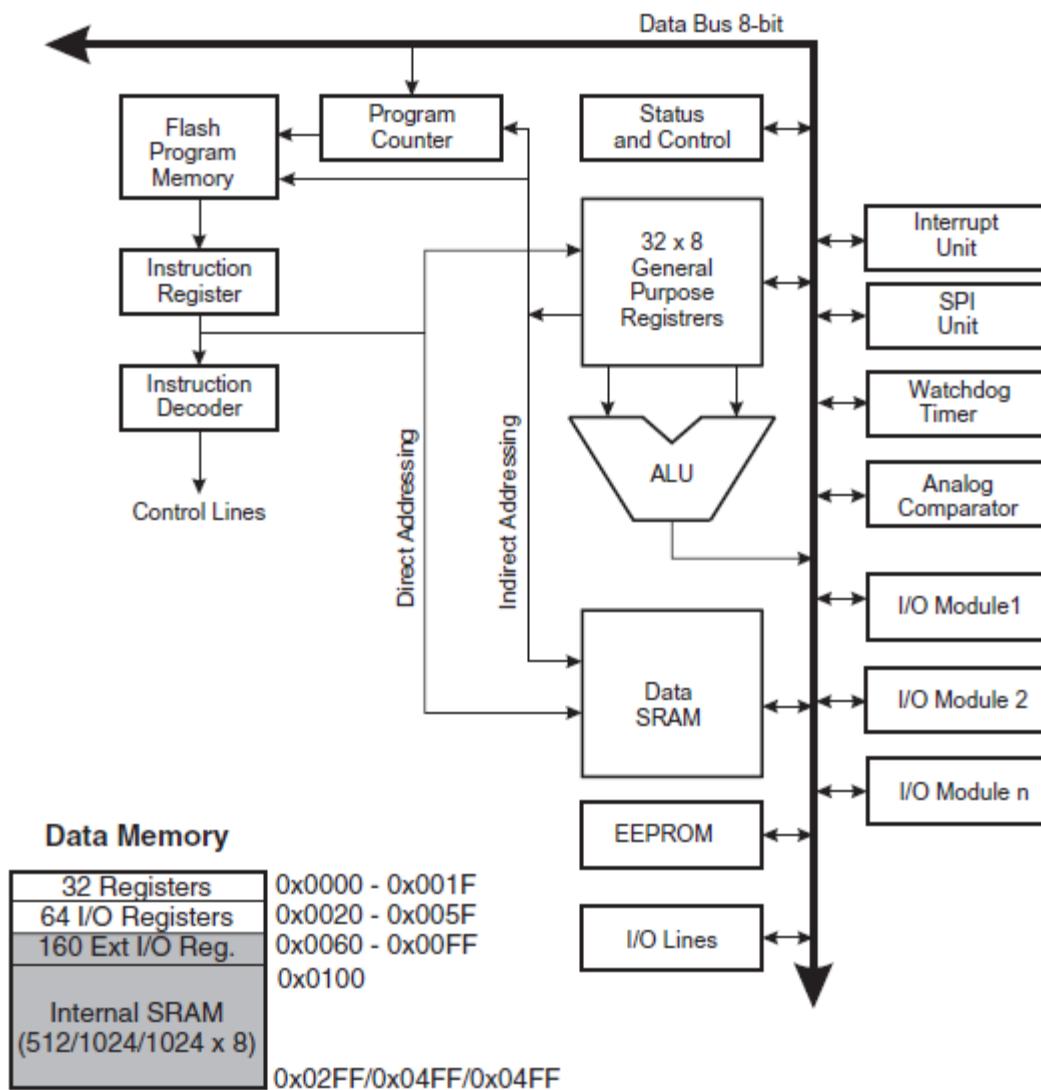
Question 3



What is meant by the following C code:

```
char *x, y;  
void foo(void) {  
    x = 0x20;  
    y = *x;  
    ...  
}
```

Answer 3




What is meant by the following C code:

```
char *x, y;  
void foo(void) {  
    x = 0x20;  
    y = *x;  
    ...  
}
```

The 8-bit quantity in the I/O register at location 0x20 is loaded into y, which is at a location in internal SRAM determined by the compiler.

Question 4



```

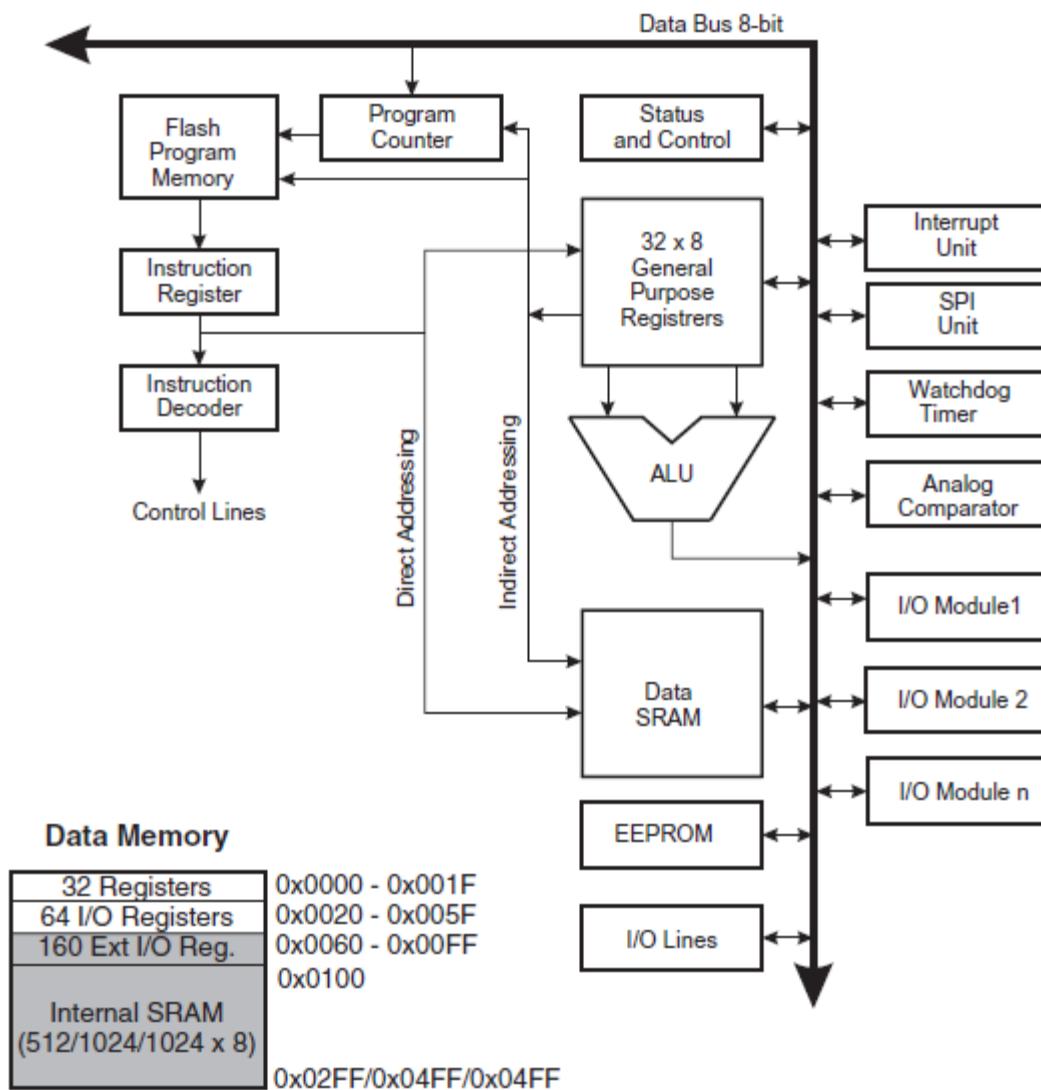
char foo() {
    char *x, y;
    x = 0x20;
    y = *x;
    return y;
}

char z;
int main(void) {
    z = foo();
    ...
}

```

Where are x, y, z in memory?

Answer 4



```


char foo() {
    char *x, y;
    x = 0x20;
    y = *x;
    return y;
}

char z;
int main(void) {
    z = foo();
    ...
}

```

x occupies 2 bytes on the stack, y occupies 1 byte on the stack, and z occupies 1 byte in static memory.

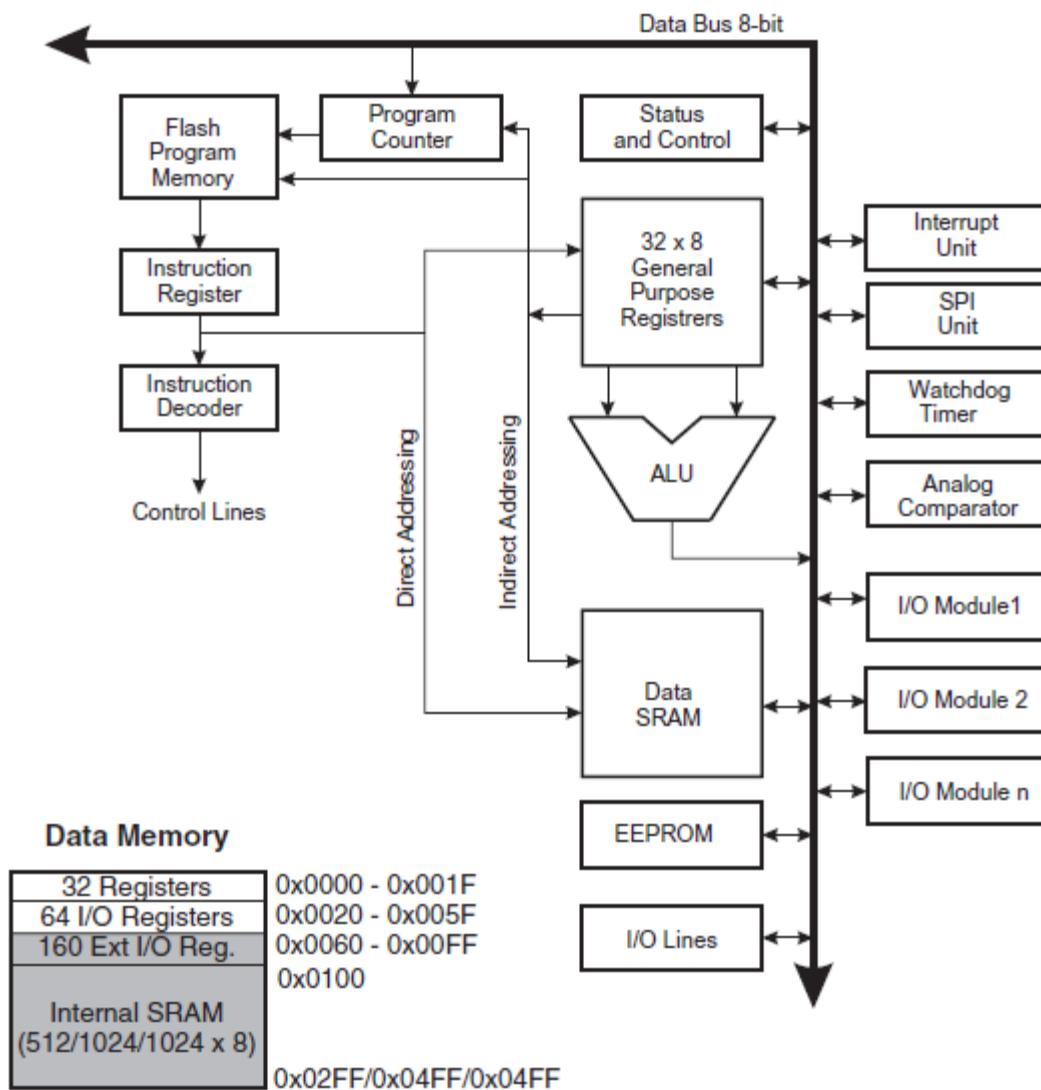
Question 5



What is meant by the following C code:

```
void foo(void) {  
    char *x, y;  
    x = &y;  
    *x = 0x20;  
    ...  
}
```

Answer 5




What is meant by the following C code:

```
void foo(void) {  
    char *x, y;  
    x = &y;  
    *x = 0x20;  
    ...  
}
```

16 bits for x and 8 bits for y are allocated on the stack, then x is loaded with the address of y, and then y is loaded with the 8-bit quantity 0x20.


Question 6

What goes into z in the following program:



```
char foo() {  
    char y;  
    uint16_t x;  
    x = 0x20;  
    y = *x;  
    return y;  
}  
  
char z;  
int main(void) {  
    z = foo();  
    ...  
}
```

Answer 6



What goes into z in the following program:

```
char foo() {  
    char y;  
    uint16_t x;  
    x = 0x20;  
    y = *x;  
    return y;  
}  
  
char z;  
int main(void) {  
    z = foo();  
    ...  
}
```

z is loaded with the 8-bit quantity in the I/O register at location 0x20.

Quiz: Find the flaw in this program (begin by thinking about where each variable is allocated)

```
int x = 2;

int* foo(int y) {
    int z;
    z = y * x;
    return &z;
}

int main(void) {
    int* result = foo(10);
    ...
}
```

Solution: Find the flaw in this program

```
int x = 2;                                statically allocated: compiler assigns a memory location.  
  
int* foo(int y){                         arguments on the stack  
    int z;                                automatic variables on the stack  
    z = y * x;  
    return &z;  
}
```

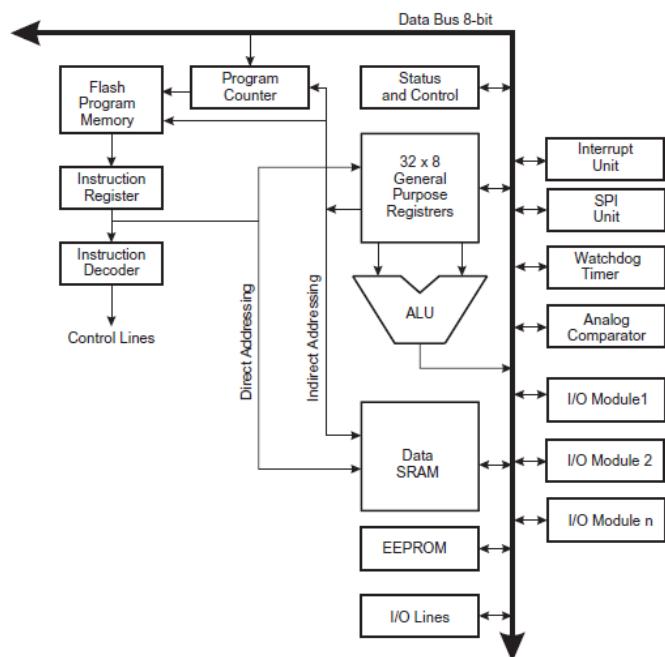
```
int main(void) {  
    int* result = foo(10);  
    ...  
}
```

program counter, argument 10,
and z go on the stack (and
possibly more, depending on the
compiler).

**The procedure foo() returns a pointer to a variable
on the stack. What if another procedure call (or
interrupt) occurs before the returned pointer is
de-referenced?**

Watch out for Recursion!!

Quiz: What is the Final Value of z?



Data Memory

32 Registers	0x0000 - 0x001F
64 I/O Registers	0x0020 - 0x005F
160 Ext I/O Reg.	0x0060 - 0x00FF
Internal SRAM (512/1024/1024 x 8)	0x0100
	...
	0x02FF/0x04FF/0x04FF

```

void foo(uint16_t x) {
    char y;
    y = *x;
    if (x > 0x100) {
        foo(x - 1);
    }
}

char z;
void main(...) {
    z = 0x10;
    foo(0x04FF);
}

```

Dynamically-Allocated Memory

The Heap

Data Memory	
32 Registers	0x0000 - 0x001F
64 I/O Registers	0x0020 - 0x005F
160 Ext I/O Reg.	0x0060 - 0x00FF
	0x0100
Internal SRAM (512/1024/1024 x 8)	0x02FF/0x04FF/0x04FF

An operating system typically offers a way to dynamically allocate memory on a “heap”.

Memory management (`malloc()` and `free()`) can lead to many problems with embedded systems:

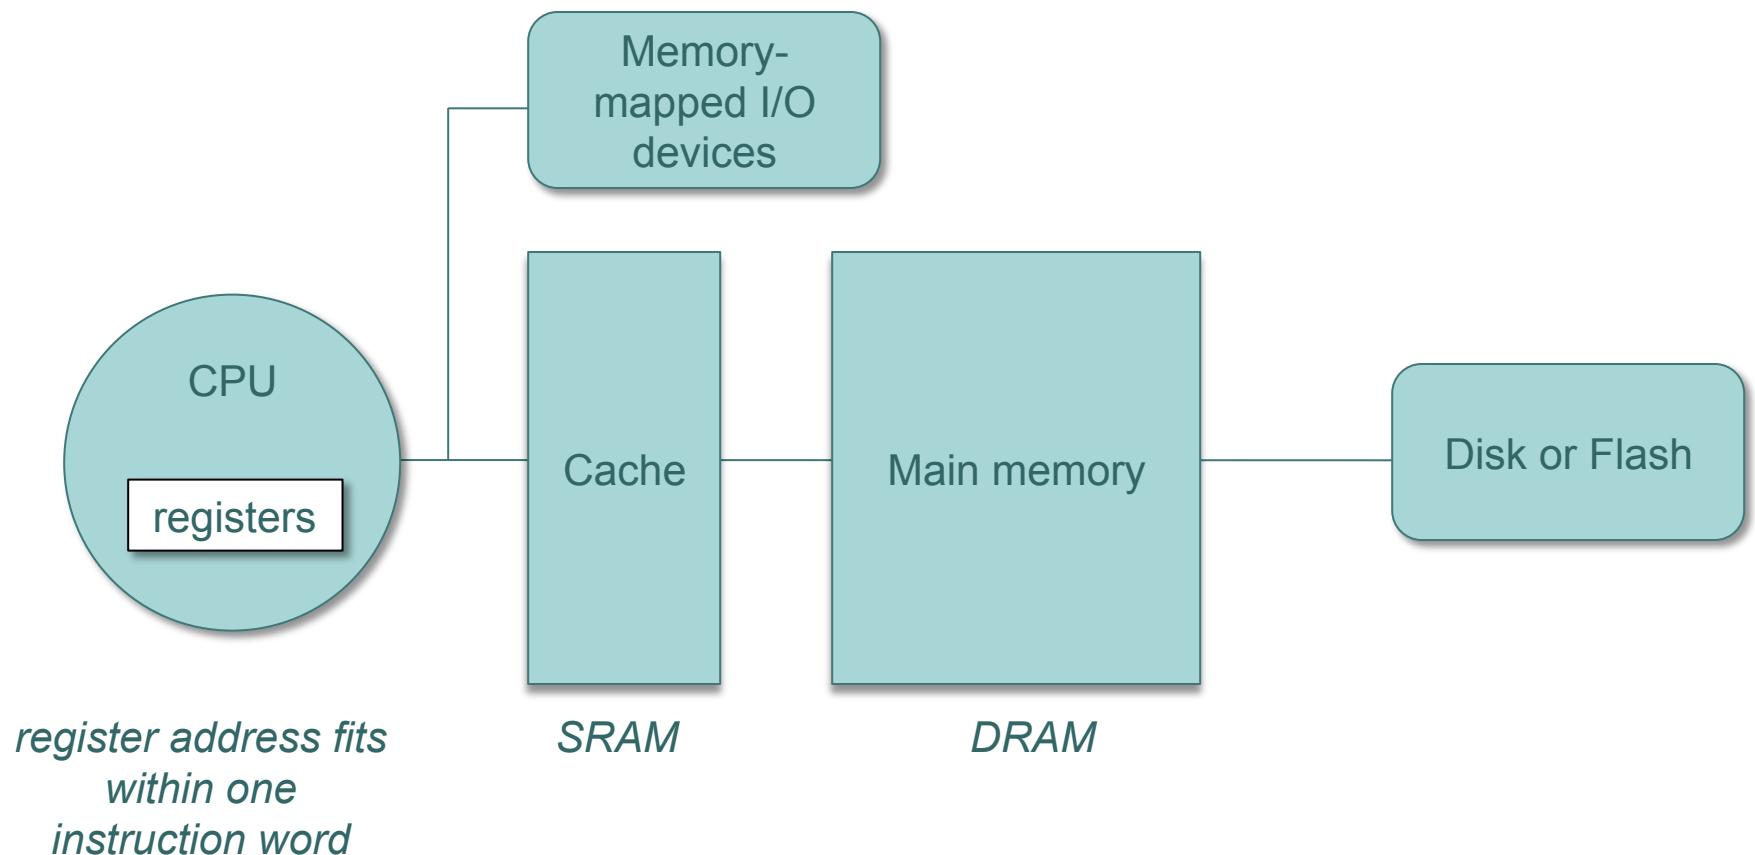
- Memory leaks (allocated memory is never freed)
- Memory fragmentation (allocatable pieces get smaller)

Automatic techniques (“garbage collection”) often require stopping everything and reorganizing the allocated memory. This is deadly for real-time programs.

Memory Hierarchies


- Memory hierarchy
 - Cache:
 - A subset of memory addresses is mapped to SRAM
 - Accessing an address not in SRAM results in *cache miss*
 - A miss is handled by copying contents of DRAM to SRAM
 - Scratchpad:
 - SRAM and DRAM occupy disjoint regions of memory space
 - Software manages what is stored where
- Segmentation
 - Logical addresses are mapped to a subset of physical addresses
 - Permissions regulate which tasks can access which memory

Memory Hierarchy



Here, the cache or scratchpad, main memory, and disk or flash share the same address space.

Memory Hierarchy




Here, each distinct piece of memory hardware has its own segment of the address space.


This requires more careful software design, but gives more direct control over timing.

Direct-Mapped Cache

A “set” consists of one “line”

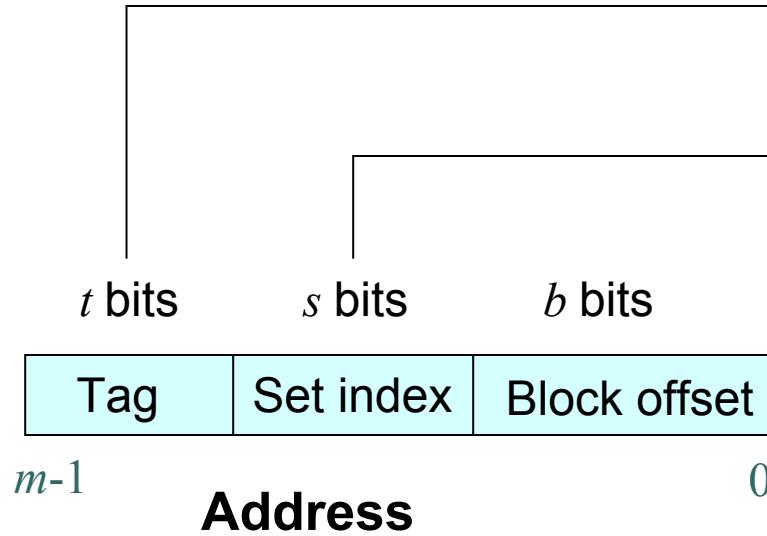


If the tag of the address matches the tag of the line, then we have a “cache hit.” Otherwise, the fetch goes to main memory, updating the line.




Set-Associative Cache

A “set” consists of several “lines”




Tag matching is done using an “associative memory” or “content-addressable memory.”



Set-Associative Cache

A “set” consists of several “lines”




A “cache miss” requires a replacement policy (like LRU or FIFO).

1 valid bit t tag bits $B = 2^b$ bytes per block


Set 0



Set 1




Set S



CACHE

Your Lab Hardware (2014 - 2016)

myRIO 1950/1900
(National Instruments)




Xilinx Zynq Z-7010

- ARM Cortex-A9 MPCore dual core processor
 - Real-time Linux
- Xilinx Artix-7 FPGA
 - Preconfigured with a 32-bit MicroBlaze microprocessor running without an operating system (“bare metal”).



Xilinx Zynq

Dual-core ARM processor + FPGA + rich I/O on a single chip.



Microblaze I/O Architecture

Source:
Xilinx




Figure 2-1: MicroBlaze Core Block Diagram

Microblaze I/O Architecture

Source:
Xilinx

M_AXI_DP: Peripheral Data Interface, AXI4-Lite or AXI4 interface

DPLB: Data interface, Processor Local Bus

DLMB: Data interface, Local Memory Bus (BRAM only)

M_AXI_IP: Peripheral Instruction interface, AXI4-Lite interface

IPLB: Instruction interface, Processor Local Bus

ILMB: Instruction interface, Local Memory Bus (BRAM only)

M0_AXIS..M15_AXIS: AXI4-Stream interface master direct connection interfaces

S0_AXIS..S15_AXIS: AXI4-Stream interface slave direct connection interfaces

MFSL 0..15: FSL master interfaces

DWFSL 0..15: FSL master direct connection interfaces

SFSL 0..15: FSL slave interfaces

DRFSL 0..15: FSL slave direct connection interfaces

DXCL: Data side Xilinx CacheLink interface (FSL master/slave pair)

M_AXI_DC: Data side cache AXI4 interface

IXCL: Instruction side Xilinx CacheLink interface (FSL master/slave pair)

M_AXI_IC: Instruction side cache AXI4 interface

Core: Miscellaneous signals for: clock, reset, debug, and trace






Figure 2-1: MicroBlaze Core Block Diagram

Berkeley Microblaze Personality Memory Map



Conclusion

Understanding memory architectures is essential to programming embedded systems.